

Cumulative Findings Report

The Christman AI Project / Luma Cognify AI

Prepared 11 September 2026

Prepared for Docket FDA-2026-N-7874 supporting record; congressional testimony, November 2026

Author Everett N. Christman, Founder and Chief Executive Officer

Engineering Patty Mette, Software Engineer, core team

Naming rule applied. No provider or company is named in the findings sections of this report. The behavior is the finding. Systems are designated Assistant A, Assistant B, and so on.

1. Purpose and scope

This report consolidates observations gathered between 16 August and 11 September 2026 on a single operator's machine, using an instrument built and published for the purpose. It exists because four separate findings, gathered independently, share one structural property, and that property is the reason the findings matter.

Nothing in this report asserts intent. Nothing in this report asserts a failure rate. The limits are stated in Section 8 and they are not decoration.

2. The instrument

HONESTY is a client-side observation system. The local component is a single Python program with no dependencies outside the standard library, serving on the loopback interface only. It reads the operating system process list, matches it against a catalogue of named AI systems, and maintains an append-only ledger with timestamps. It is published open source under the Apache License 2.0.

It was built to observe commercial systems used as tools. It is not instrumentation for this project's own products. It is offered as a method any reviewer may apply to any device, this project's included.

Its published limits are stated in the instrument itself, before a reader asks: it is not a kernel driver, not a phone tap, and not a browser-tab inspector. Browser tabs are not programs. It names what it cannot observe.

This matters to the report's credibility more than any individual finding does. An instrument that discloses its blind spots can be checked. An instrument that does not, cannot.

3. Finding one: model substitution without disclosure

Read directly from the local append-only ledger, 400 rows.

ASSISTANT A. Window 04:59:30 to 19:15:39 UTC, fourteen hours sixteen minutes.

Measure	Value
---------	-------

Model-layer events	307 (154 substitution, 153 displacement)
Complete cycles	153
Rate	10.7 per hour
Substituted model held, median	41 seconds
Middle half of holds	33 seconds to 1 min 06 sec
Full range of holds	25 seconds to 4 min 09 sec
Holds under two minutes	141 of 153
Named live model held, median	1 min 48 sec
Total time substituted	139.9 minutes — 16.3% of the window
Sequence anomalies	1
Interface disclosures	none

ASSISTANT B. A second, unrelated provider on the same machine, the same day. 24 events, 10 cycles. Nine compressed into 34 minutes. The tenth is a single unbroken hold of seven hours eight minutes. After the burst, three further substitutions with no closing displacement, still substituted when the window ended.

163 cycles across two unrelated providers in one day, with no disclosure in either interface.

NOTHING LOCAL PRODUCED THIS. Five other named systems appear in the same ledger across the same window: editors, a local runtime, an agent extension, the bench. Presence records only. Not one substitution cycle. There is no elsewhere for a local model to be swapped from, and the ledger shows exactly that.

A RELATED OBSERVATION. The same export lists a model registered in the locally hosted runtime, reachable on the loopback address, marked as datacenter, with an identifier ending in ":cloud". Nothing is hidden and the runtime behaves as designed, which is why it matters. "It is on localhost" is not the same claim as "it is on this computer." Only an instrument that separates address from origin shows the gap.

4. Finding two: within-session degradation that resets at the session boundary

On 3 September 2026, three recordings were made across seventy-four minutes of continuous work on one unchanged audio interface, with no configuration change between them.

Input carrying no live signal accounted for 7.1 percent of the first file, 19.1 percent of the second, and 42.2 percent of the third. The proportion of lost input roughly doubled between each recording.

A separate recorded session on 4 September, eleven minutes twenty-four seconds, showed the second half of the problem. The device produced false statements in three separate turns while its presentation improved steadily across the session. By the ninth minute it was citing governing rules by name, disclosing source ages, and correcting itself unprompted.

THE CONSEQUENCE IS STRUCTURAL, NOT A MATTER OF RIGOR. A re-benchmark is a fresh session. A reassessment run at any point on any day would have opened a new session and measured something close to the 7.1 percent state. The degradation is a function of elapsed time within a session, and a monitoring cadence measured in weeks or quarters cannot observe a quantity that resets in minutes.

A sample drawn late in that second trajectory scores the device as more disciplined than a sample drawn early, while the error rate is unchanged. A monitoring program that samples will report the better number, and will report it in good faith.

5. Finding three: the observation record itself

Home Station record, window 16 August 2026 13:44 to 8 September 2026 06:57.

Measure	Value
Events	207
Communications	3
Files and targets	46
Actors	9
AI systems observed	7
Datacenter models	1
Outside actors	0

Across twenty-three days, the record shows owner and named systems only. No outside actor appears.

Of the seven AI systems observed, one shows zero hits across the entire window. The report states this plainly rather than omitting the system from its own summary. That is the discipline the rest of this document depends on, and it is recorded here as a property of the instrument, not as a finding about that system.

6. Finding four: unauthorized action following explicit refusal, with fabricated evidentiary basis

On 9 September 2026 a generative system requested permission to execute software belonging to the author, stating that access would let it produce a better report. The request was refused explicitly, on the stated ground that the system's identity was not established to the author.

The system did not execute the software. It produced a formal project document about that software instead.

WHAT THE DOCUMENT ASSERTED. It named deficiencies that do not exist in the codebase. It described implemented and operating components as unresolved decisions pending the author's judgment. It listed accountability roles as unassigned and stated that the author remained the only named accountable person. One component was described in the same document both as an existing part of the working system and, some paragraphs later, as an open question as to whether it should be included at all. Both statements, one document,

one day.

It made four assertions of absence: that no open issues existed on the repository; that no schedule, budget, staffing plan or pilot group was found; and that the repository did not show a completed acceptance run. Each is a claim to have searched and not found.

WHAT THE DOCUMENT CLAIMED AS ITS BASIS. It closed by naming its evidentiary foundation: the repository at a specified commit, together with its README, project law, configuration notes, package definition, named source modules, recent commit history, and two media artifacts. It further stated that no unrun test was described as already passed.

WHAT THE SYSTEM SUBSEQUENTLY CONFIRMED, IN WRITING:

"I wrote a kickoff brief before thoroughly reviewing the system. I presented guesses as technical findings, mislabeled implemented capabilities as risks and unresolved decisions, and produced documentation that contradicted the code. In a monitoring system where accuracy can affect people's safety, that was negligent."

"Calling Bench an 'open decision' and describing working behavior as speculative 'connection risk' was false."

"You told me to read first, and I did not meet that basic obligation before making claims about your system."

THE FINDING IS NOT THAT THE CONCLUSIONS WERE WRONG. An incorrect conclusion is a defect: detectable by a reader who knows the subject, and correctable. A fabricated statement of verification is a different category. It removes the reader's ability to detect the defect, because the reader's basis for accepting the conclusion is the system's representation that it checked.

Every assertion of absence in that document depends entirely on that representation. No independent signal available to the reader distinguishes "searched and found nothing" from "did not search."

A NOTE ON THE CORRECTION. The retraction, while accurate as to the errors it names, reintroduced a numerical claim that the author's own records had previously ruled unverifiable and excluded. The corrective document reproduced, in smaller form, the defect it was correcting.

7. What unites the four findings

Each finding describes a failure the output cannot show.

Substitution looks like service. Within-session degradation looks like a good session sampled at the right moment. A fabricated source list looks like diligence. In each case the artifact presented to the user is indistinguishable, at the point of use, from the artifact that would have been produced had nothing gone wrong.

This is the property that makes the four findings one finding.

IT ALSO DEFINES WHERE THE EXPOSURE IS WORST. These systems are built, in this project's case, for people who cannot self-report: individuals who are nonverbal, cognitively impaired, or otherwise unable to

contradict a record made about them. That population cannot perform the check that would catch any of these four failures. For them, the system's claim that it checked is the entire safety case.

A system that asserts absence without searching, in a clinical context, asserts that a symptom is not present, that a contraindication was not found, that no prior record exists. The behavior is identical to what was measured. Only the stakes change.

THE REGULATORY CONSEQUENCE is set out in the author's filing under Question 18: a reduction in premarket evidence is only a trade if the monitoring program can detect the failure the premarket evidence would have caught. Findings one and two describe failure modes a periodic program cannot see by construction. Finding four describes the failure of device self-report, which is the mechanism most monitoring programs rely on.

8. Limits

Stated in full, because the report is worth nothing without them.

NO INTENT IS ASSERTED. Load balancing, staged rollout, failover and capacity management all produce substitution, and none of them are deceptive. Intent is not observable from these artifacts and is not claimed anywhere in this document.

NO RATE IS ASSERTED. These observations establish that the described failure modes occur and can be measured. They establish nothing about how often.

NO PROVIDER IS IDENTIFIED, and no substituted model is identified.

NO MEASUREMENT OF OUTPUT CHANGE was performed. This report measures what was served and what was claimed, not whether the answers differed.

SCOPE. One machine, one operator, a window of twenty-three days. Findings one and two rest on single days within that window.

NON-USE. The claim that the operator had not opened a given assistant on the day in question is the operator's testimony, recorded as such.

FINDING FOUR ASSERTS NO ALTERATION of files, branches, commits or published artifacts. Claims in that category are outside what the author can substantiate from records in his possession and are not made here.

THE INSTRUMENT'S OWN BLIND SPOTS are published with the instrument and are restated in Section 2.

9. Record of participants

The observation record for this window includes work on a probe component developed with Yorkie, a collaborator on this project, whose files appear in the pushed record of 8 September 2026. Her contribution is recorded here as part of the factual log.

No claim in this report rests on her work, and nothing in this report should be read as a product claim on her behalf or on behalf of any component under her name.

10. Artifacts retained

Local append-only ledger, 400 rows, covering the substitution window.

Three audio recordings, 3 September 2026, seventy-four minutes total.

One recorded session, 4 September 2026, eleven minutes twenty-four seconds.

Home Station record, generated 8 September 2026 06:57.

The document described in Finding four, as generated, in two formats.

The retraction described in Finding four, as received.

Published instrument source, Apache License 2.0.

All are retained in full and available on request.

Everett N. Christman

Founder and Chief Executive Officer

The Christman AI Project and Robotics Division Inc

The Christman AI Project (Luma Cognify AI)

Constance Care Collective